

```
In [1]:
```

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.tree import DecisionTreeClassifier, plot_tree
from mlxtend.plotting import plot_decision_regions
from sklearn.preprocessing import LabelEncoder
import warnings
warnings.filterwarnings("ignore", category=pd.errors.SettingWithCopyWarning)
```

```
In [3]:
```

```
df = pd.read_csv(r"C:\Users\Admin\Downloads\archive\recipes_muffins_cupcakes.csv") # change
df.head()
```

```
Out[3]:
```

	Type	Flour	Milk	Sugar	Butter	Egg	Baking Powder	Vanilla	Salt	
0	Muffin	55	28	3	7	5		2	0	0
1	Muffin	47	24	12	6	9		1	0	0
2	Muffin	47	23	18	6	4		1	0	0
3	Muffin	45	11	17	17	8		1	0	0
4	Muffin	50	25	12	6	5		2	1	0

```
In [4]:
```

```
df.shape
```

```
Out[4]:
```

```
(20, 9)
```

```
In [5]:
```

```
le = LabelEncoder()
df['label'] = le.fit_transform(df['Type'])
df['label']
```

```
Out[5]:
```

```
0      1
1      1
2      1
3      1
4      1
5      1
6      1
7      1
8      1
9      1
10     0
11     0
12     0
13     0
14     0
15     0
16     0
17     0
18     0
19     0
Name: label, dtype: int64
```

```
In [6]:
```

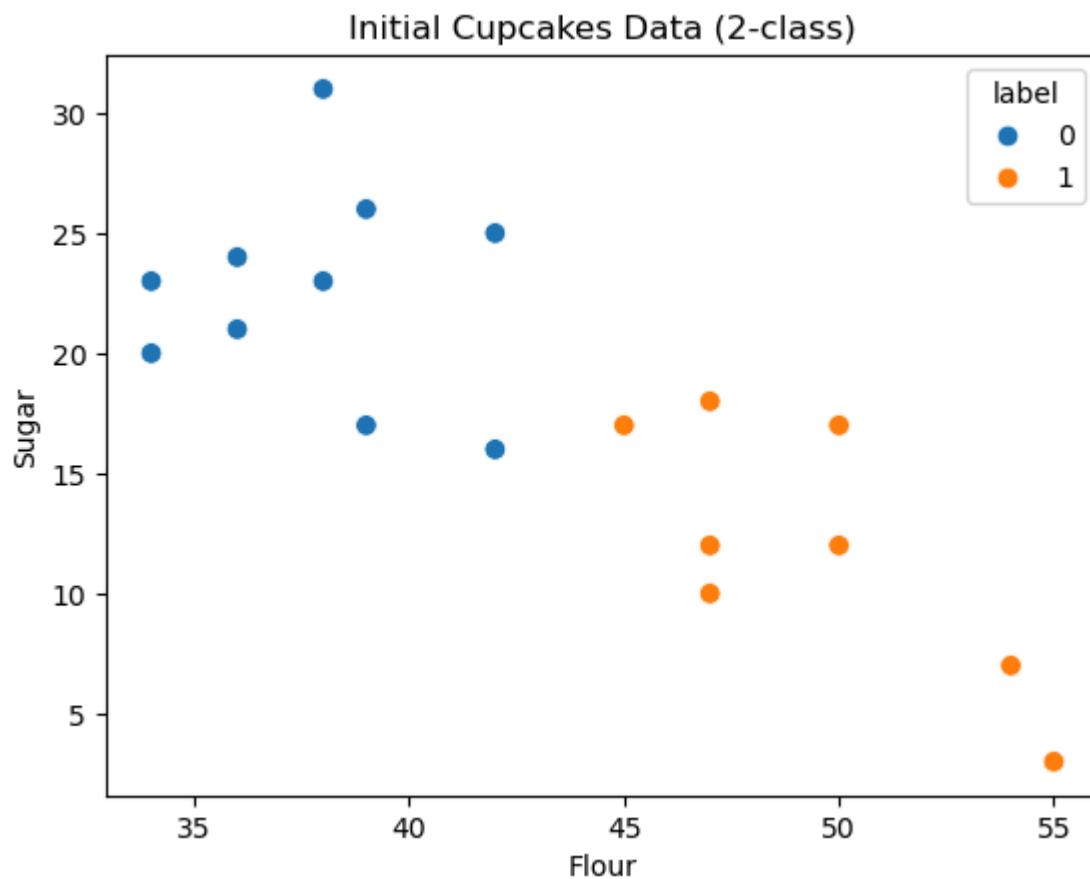
```
df = df[['Flour', 'Sugar', 'label']]
df.columns = ["X1", "X2", "label"]
```

```
In [7]:
```

```
df["weights"] = 1 / df.shape[0]
display(df.head())
```

	X1	X2	label	weights
0	55	3	1	0.05
1	47	12	1	0.05
2	47	18	1	0.05
3	45	17	1	0.05
4	50	12	1	0.05

```
In [8]: sns.scatterplot(x="X1", y="X2", hue="label", data=df, s=60)
plt.title("Initial Cupcakes Data (2-class)")
plt.xlabel("Flour")
plt.ylabel("Sugar")
plt.show()
```



```
In [9]: x = df[["X1", "X2"]].values
y = df["label"].values
```

```
In [10]: def calculate_model_weight(error):
    eps = 1e-10
    error = np.clip(error, eps, 1 - eps)
    return 0.5 * np.log((1 - error) / error)
```

```
In [11]: def update_row_weights(row, alpha):
    if row["label"] == row["y_pred"]:
        return row["weights"] * np.exp(-alpha)
    else:
        return row["weights"] * np.exp(alpha)

def create_new_dataset(df):
    indices = []
    n = df.shape[0]
```

```

for _ in range(n):
    a = np.random.random()
    for idx, row in df.iterrows():
        if row["cumsum_lower"] < a <= row["cumsum_upper"]:
            indices.append(idx)
            break
    if len(indices) == 0:
        indices = np.random.choice(df.index, size=n, replace=True).tolist()
return indices

```

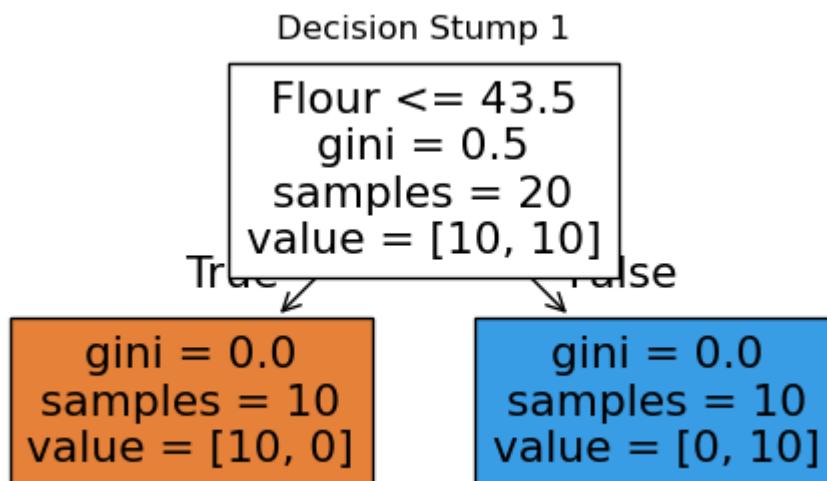
In [12]: #Step 2: Train First Decision Stump  
`dt1 = DecisionTreeClassifier(max_depth=1)  
dt1.fit(x, y)`

```

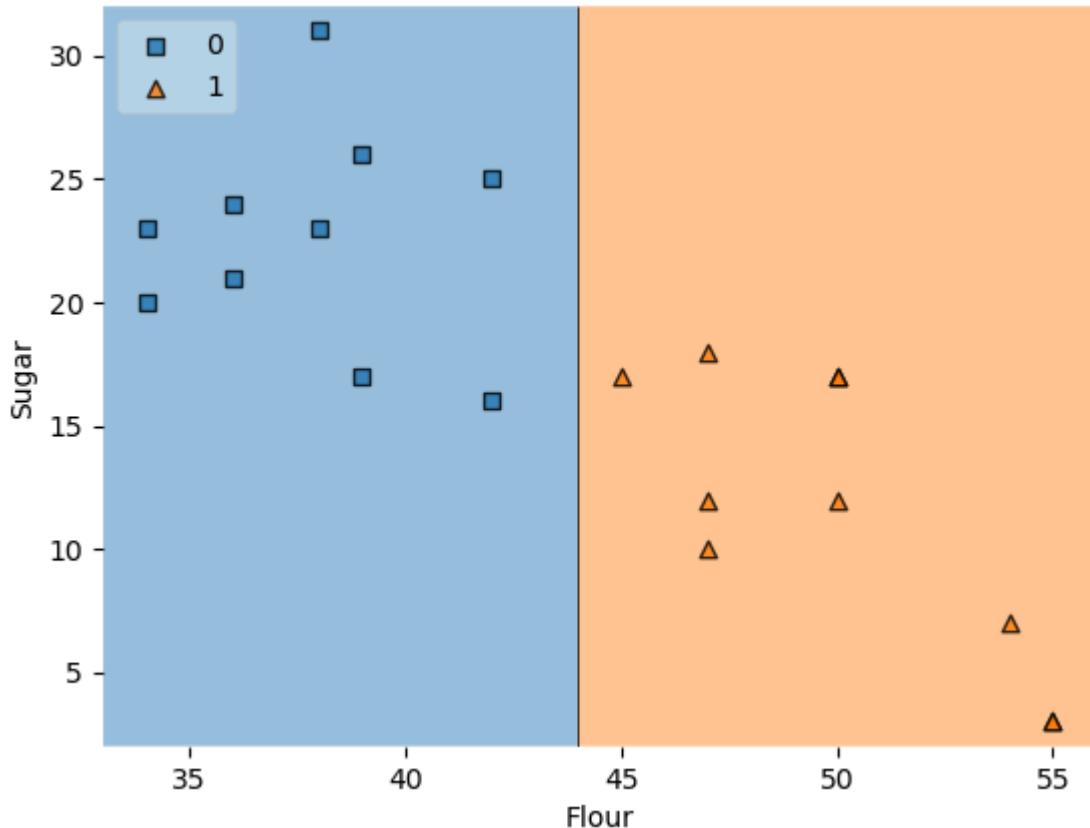
plt.figure(figsize=(6,3))
plot_tree(dt1, filled=True, feature_names=["Flour", "Sugar"])
plt.title("Decision Stump 1")
plt.show()

plot_decision_regions(x, y, clf=dt1, legend=2)
plt.title("Decision Region - Stump 1")
plt.xlabel("Flour")
plt.ylabel("Sugar")
plt.show()

```



## Decision Region - Stump 1



```
In [13]: # Predictions & weighted error
df["y_pred"] = dt1.predict(x)
error = np.sum(df["weights"] * (df["label"] != df["y_pred"]))
alpha1 = calculate_model_weight(error)
print(f"Model 1 alpha: {alpha1:.3f}")

# Update and normalize weights
df["updated_weights"] = df.apply(lambda row: update_row_weights(row, alpha1), axis=1)
df["normalized_weights"] = df["updated_weights"] / df["updated_weights"].sum()

# Compute cumulative bounds
df["cumsum_upper"] = np.cumsum(df["normalized_weights"])
df["cumsum_lower"] = df["cumsum_upper"] - df["normalized_weights"]
display(df[["X1", "X2", "label", "weights", "y_pred", "normalized_weights"]])
```

Model 1 alpha: 11.513

	X1	X2	label	weights	y_pred	normalized_weights
0	55	3	1	0.05	1	0.05
1	47	12	1	0.05	1	0.05
2	47	18	1	0.05	1	0.05
3	45	17	1	0.05	1	0.05
4	50	12	1	0.05	1	0.05
5	55	3	1	0.05	1	0.05
6	54	7	1	0.05	1	0.05
7	47	10	1	0.05	1	0.05
8	50	17	1	0.05	1	0.05
9	50	17	1	0.05	1	0.05
10	39	26	0	0.05	0	0.05
11	42	16	0	0.05	0	0.05
12	34	20	0	0.05	0	0.05
13	39	17	0	0.05	0	0.05
14	38	23	0	0.05	0	0.05
15	42	25	0	0.05	0	0.05
16	36	21	0	0.05	0	0.05
17	38	31	0	0.05	0	0.05
18	36	24	0	0.05	0	0.05
19	34	23	0	0.05	0	0.05

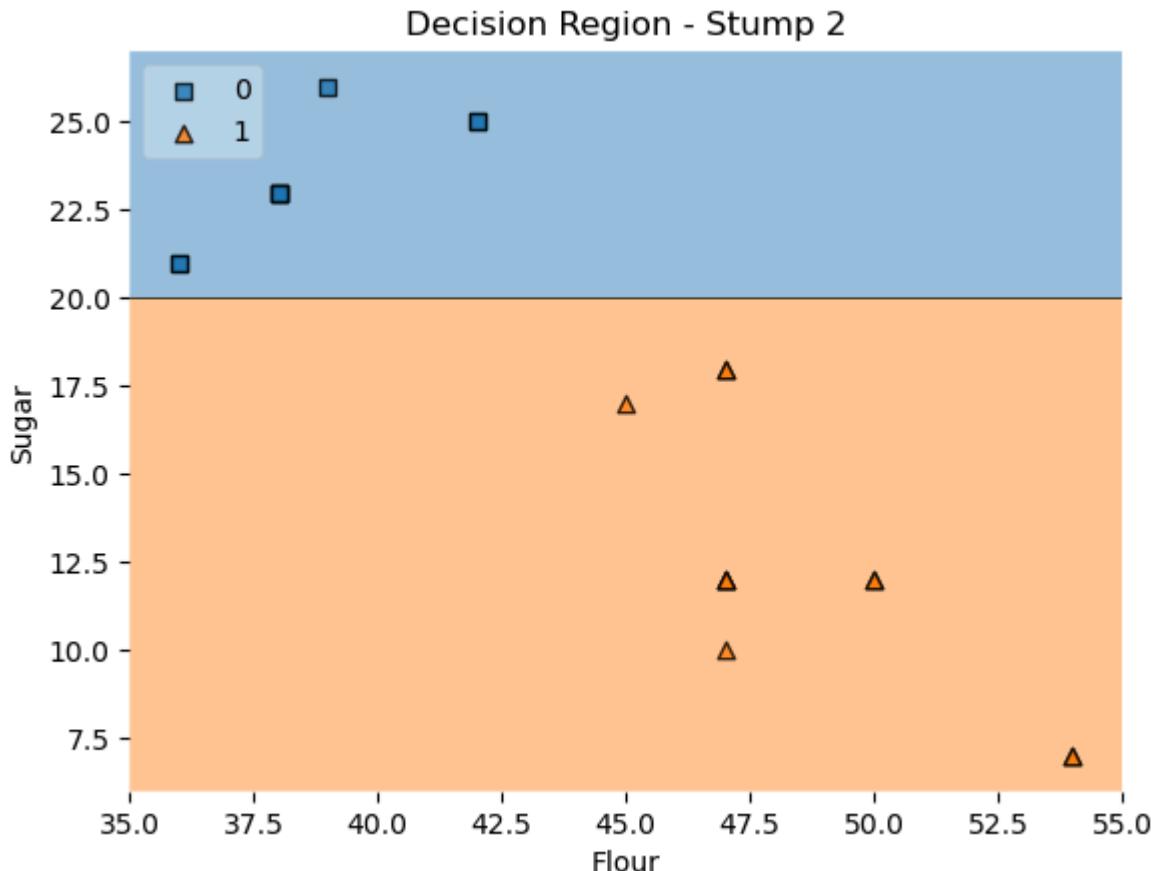
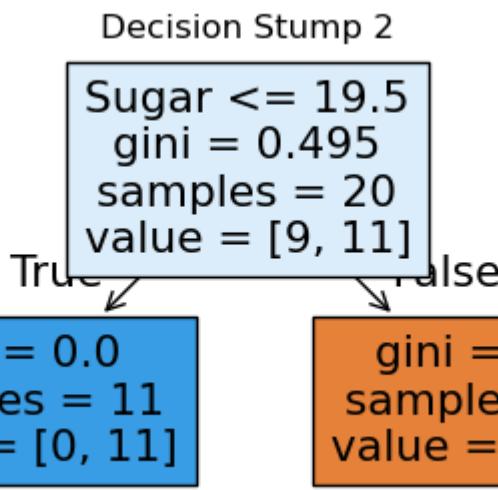
```
In [16]: # Resample for next stump
index_values = create_new_dataset(df)
print("Resampled indices (1st):", index_values)
second_df = df.loc[index_values, ["X1", "X2", "label", "normalized_weights"]].reset_index(drop=True)
```

Resampled indices (1st): [16, 14, 15, 4, 1, 2, 14, 7, 10, 14, 6, 1, 16, 2, 6, 1, 4, 15, 3, 14]

```
In [17]: # =====
# Step 3: Train Second Decision Stump
# =====
x2 = second_df[["X1", "X2"]].values
y2 = second_df["label"].values
dt2 = DecisionTreeClassifier(max_depth=1)
dt2.fit(x2, y2)

plt.figure(figsize=(6,3))
plot_tree(dt2, filled=True, feature_names=["Flour", "Sugar"])
plt.title("Decision Stump 2")
plt.show()

plot_decision_regions(x2, y2, clf=dt2, legend=2)
plt.title("Decision Region - Stump 2")
plt.xlabel("Flour")
plt.ylabel("Sugar")
plt.show()
```



```
In [18]: second_df["y_pred"] = dt2.predict(x2)
error2 = np.sum(second_df["normalized_weights"] * (second_df["label"] != second_df["y_pred"]))
alpha2 = calculate_model_weight(error2)
print(f"Model 2 alpha: {alpha2:.3f}")
```

Model 2 alpha: 11.513

```
In [19]: # Update weights again
def update_row_weights_2(row, alpha=alpha2):
    if row["label"] == row["y_pred"]:
        return row["normalized_weights"] * np.exp(-alpha)
    else:
        return row["normalized_weights"] * np.exp(alpha)

second_df["updated_weights"] = second_df.apply(update_row_weights_2, axis=1)
second_df["normalized_weights"] = second_df["updated_weights"] / second_df["updated_weights"]
second_df["cumsum_upper"] = np.cumsum(second_df["normalized_weights"])
second_df["cumsum_lower"] = second_df["cumsum_upper"] - second_df["normalized_weights"]

index_values2 = create_new_dataset(second_df)
```

```
print("Resampled indices (2nd):", index_values2)
third_df = second_df.loc[index_values2, ["X1", "X2", "label", "normalized_weights"]].reset_index()
```

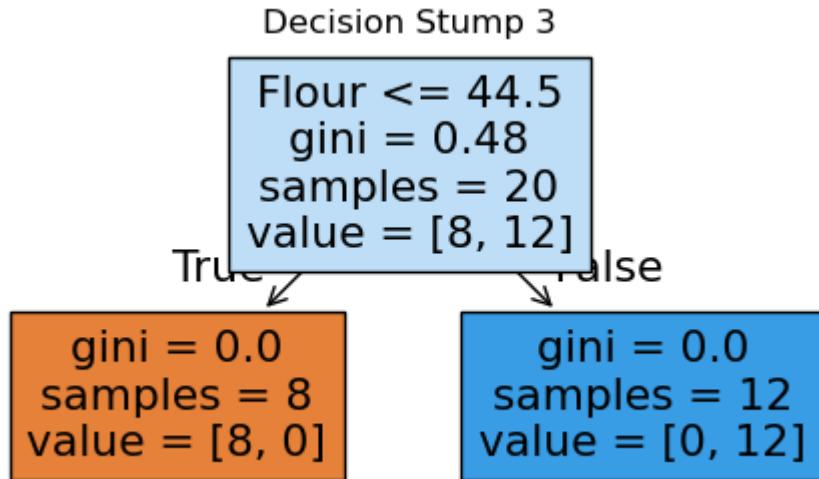
```
Resampled indices (2nd): [16, 16, 3, 17, 5, 8, 2, 7, 11, 16, 9, 0, 5, 11, 19, 16, 8, 3, 13, 19]
```

```
In [20]: # =====
```

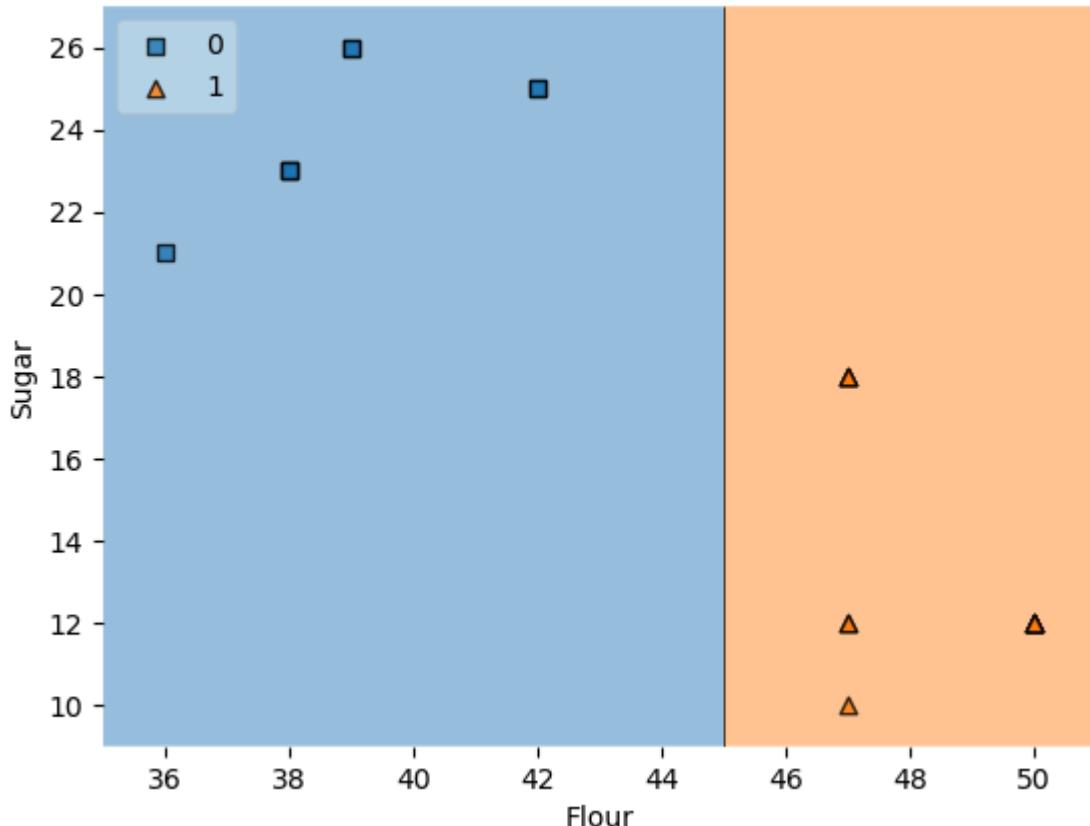
```
# Step 4: Train Third Decision Stump
# =====
x3 = third_df[["X1", "X2"]].values
y3 = third_df["label"].values
dt3 = DecisionTreeClassifier(max_depth=1)
dt3.fit(x3, y3)

plt.figure(figsize=(6,3))
plot_tree(dt3, filled=True, feature_names=["Flour", "Sugar"])
plt.title("Decision Stump 3")
plt.show()

plot_decision_regions(x3, y3, clf=dt3, legend=2)
plt.title("Decision Region - Stump 3")
plt.xlabel("Flour")
plt.ylabel("Sugar")
plt.show()
```



### Decision Region - Stump 3



```
In [21]: third_df["y_pred"] = dt3.predict(x3)
error3 = np.sum(third_df["normalized_weights"] * (third_df["label"] != third_df["y_pred"]))
alpha3 = calculate_model_weight(error3)
print(f"Model 3 alpha: {alpha3:.3f}")

Model 3 alpha: 11.513
```

```
In [22]: # Step 5: Combined Visualization
# =====
print("\nFinal Model Weights (alphas):")
print(f"Alpha1 = {alpha1:.3f}, Alpha2 = {alpha2:.3f}, Alpha3 = {alpha3:.3f}")

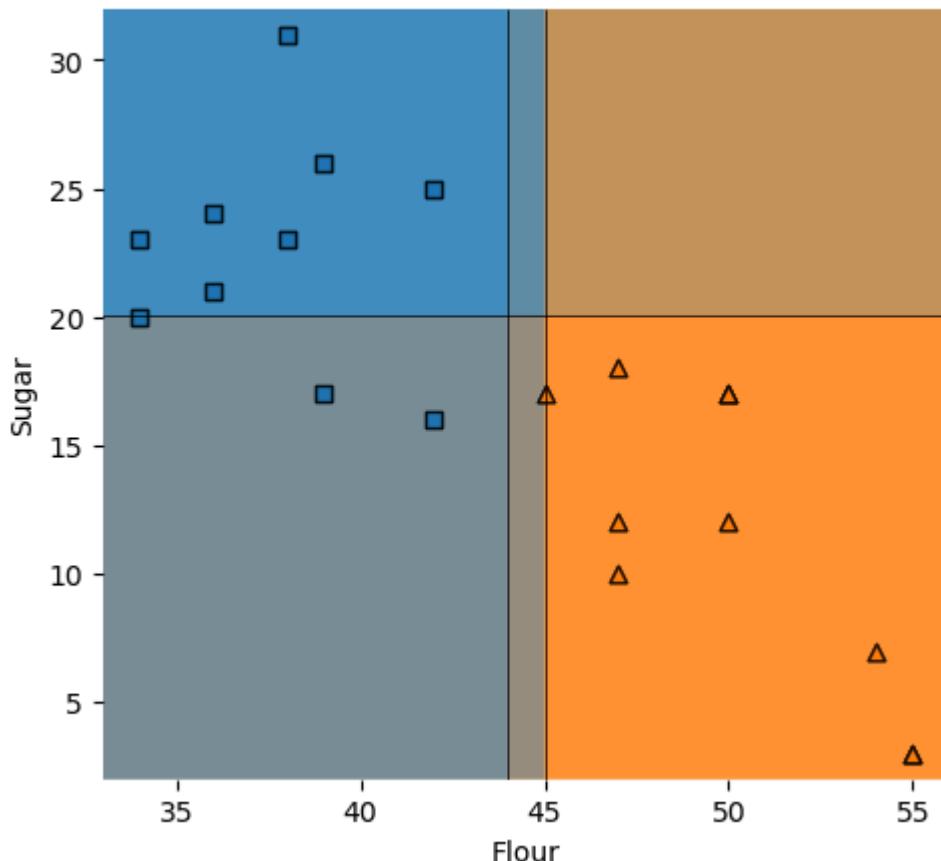
plt.figure(figsize=(12,5))

plt.subplot(1,2,1)
plot_decision_regions(x, y, clf=dt1, legend=0)
plot_decision_regions(x, y, clf=dt2, legend=0)
plot_decision_regions(x, y, clf=dt3, legend=0)
plt.title("Decision Regions: All 3 Stumps")
plt.xlabel("Flour")
plt.ylabel("Sugar")
```

Final Model Weights (alphas):  
Alpha1 = 11.513, Alpha2 = 11.513, Alpha3 = 11.513

Out[22]: Text(0, 0.5, 'Sugar')

### Decision Regions: All 3 Stumps

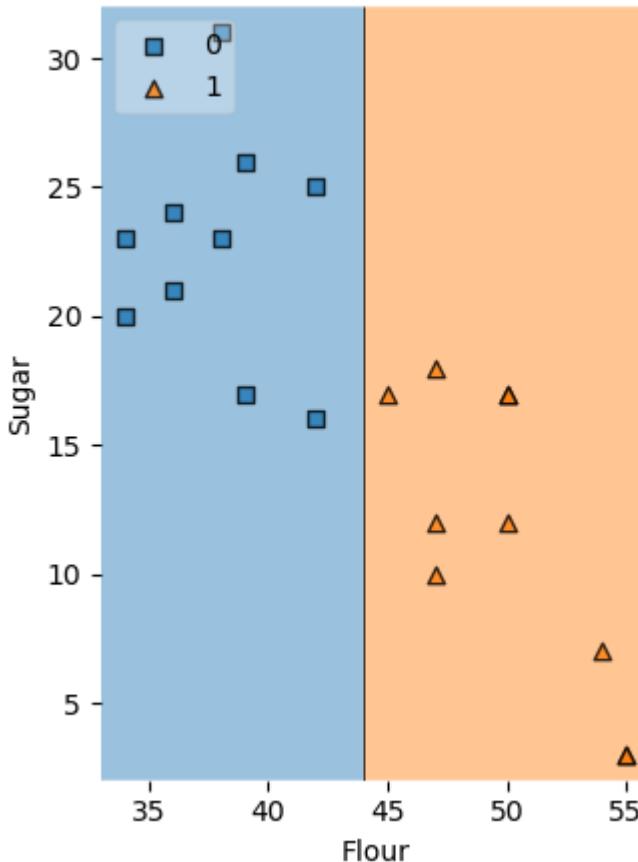


```
In [23]: from sklearn.ensemble import AdaBoostClassifier
ada_final = AdaBoostClassifier(
    estimator=DecisionTreeClassifier(max_depth=1),
    n_estimators=3,
    algorithm="SAMME"
)
ada_final.fit(x, y)

plt.subplot(1,2,2)
plot_decision_regions(x, y, clf=ada_final, legend=2)
plt.title("Final AdaBoost Combined Decision Region")
plt.xlabel("Flour")
plt.ylabel("Sugar")
plt.tight_layout()
plt.show()
```

C:\Users\Admin\.conda\envs\myenv\Lib\site-packages\sklearn\ensemble\\_weight\_boosting.py:519: FutureWarning: The parameter 'algorithm' is deprecated in 1.6 and has no effect. It will be removed in version 1.8.  
warnings.warn(

## Final AdaBoost Combined Decision Region



```
In [24]: # =====
```

```
# Step 6: Visualize All Stumps Together
# =====
from sklearn.ensemble import AdaBoostClassifier

plt.figure(figsize=(14, 6))

# --- Subplot 1: Decision Stump 1 ---
plt.subplot(1, 4, 1)
plot_decision_regions(x, y, clf=dt1, legend=0)
plt.title("Decision Stump 1")
plt.xlabel("X1")
plt.ylabel("X2")

# --- Subplot 2: Decision Stump 2 ---
plt.subplot(1, 4, 2)
plot_decision_regions(x2, y2, clf=dt2, legend=0)
plt.title("Decision Stump 2")
plt.xlabel("X1")
plt.ylabel("X2")

# --- Subplot 3: Decision Stump 3 ---
plt.subplot(1, 4, 3)
plot_decision_regions(x3, y3, clf=dt3, legend=0)
plt.title("Decision Stump 3")
plt.xlabel("X1")
plt.ylabel("X2")

# --- Subplot 4: Final Combined AdaBoost ---
ada_final = AdaBoostClassifier(
    estimator=DecisionTreeClassifier(max_depth=1),
    n_estimators=3,
    algorithm="SAMME"
)
ada_final.fit(x, y)
```

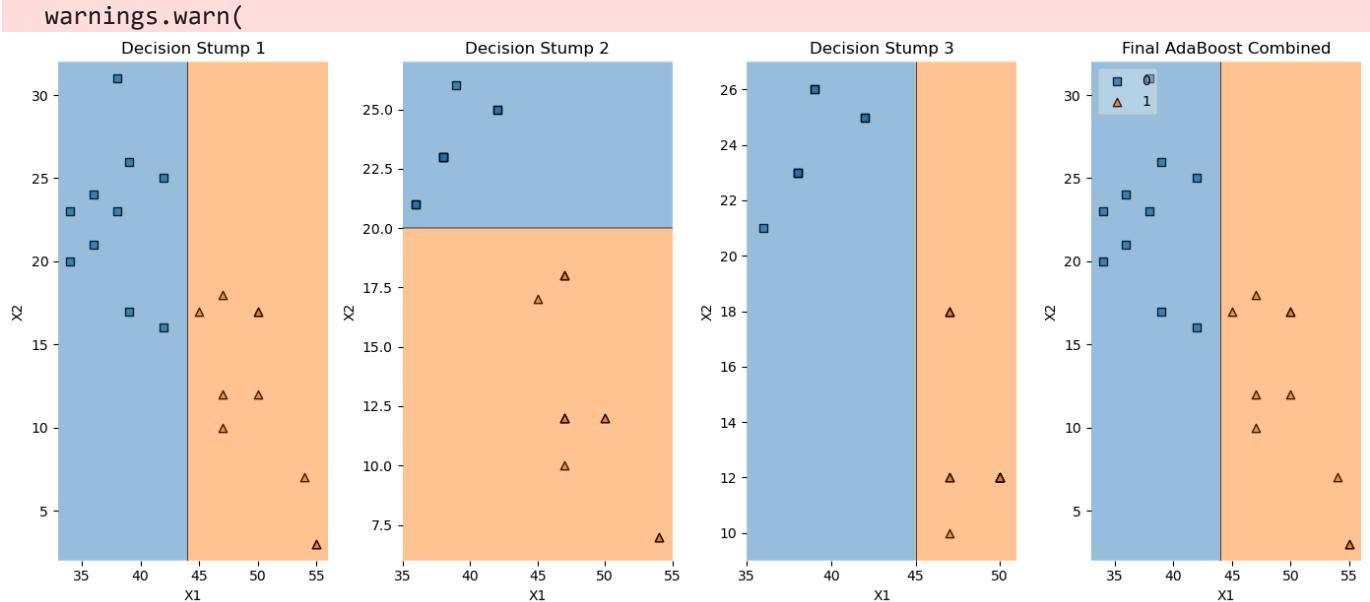
```

plt.subplot(1, 4, 4)
plot_decision_regions(x, y, clf=ada_final, legend=2)
plt.title("Final AdaBoost Combined")
plt.xlabel("X1")
plt.ylabel("X2")

plt.tight_layout()
plt.show()

```

C:\Users\Admin\.conda\envs\myenv\Lib\site-packages\sklearn\ensemble\\_weight\_boosting.py:519: FutureWarning: The parameter 'algorithm' is deprecated in 1.6 and has no effect. It will be removed in version 1.8.



In [25]:

```

# -----
# 🍩 Muffin vs Cupcake using AdaBoost + ALL Decision Stumps Visualization
# -----


import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.ensemble import AdaBoostClassifier

```

In [27]:

```
df = pd.read_csv(r"C:\Users\Admin\Downloads\archive\recipes_muffins_cupcakes.csv") # change
df.head()
```

Out[27]:

	Type	Flour	Milk	Sugar	Butter	Egg	Baking Powder	Vanilla	Salt
0	Muffin	55	28	3	7	5	2	0	0
1	Muffin	47	24	12	6	9	1	0	0
2	Muffin	47	23	18	6	4	1	0	0
3	Muffin	45	11	17	17	8	1	0	0
4	Muffin	50	25	12	6	5	2	1	0

In [28]:

```

le = LabelEncoder()
df['label'] = le.fit_transform(df['Type'])
df['label']

```

```
Out[28]: 0      1
1      1
2      1
3      1
4      1
5      1
6      1
7      1
8      1
9      1
10     0
11     0
12     0
13     0
14     0
15     0
16     0
17     0
18     0
19     0
Name: label, dtype: int64
```

```
In [29]: df['Label'] = df['Type'].map({'Muffin': 0, 'Cupcake': 1})
```

```
In [30]: X = df[['Sugar', 'Flour']]
y = df['Label']
```

```
In [31]: # 2 Split into training and test data
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42)
```

```
In [32]: # 3 Create AdaBoost model with Decision Stumps
base_tree = DecisionTreeClassifier(max_depth=1)
ada = AdaBoostClassifier(
    estimator=base_tree,
    n_estimators=3,          # use 3 stumps for visualization
    learning_rate=1.0,
    random_state=42
)
```

```
In [34]: # 4 Train model
ada.fit(X_train, y_train)
```

```
Out[34]: ▶   AdaBoostClassifier
           ⓘ ⓘ
           ▶     estimator:
           DecisionTreeClassifier
           ▶       DecisionTreeClassifier ⓘ
```

```
In [35]: from sklearn.tree import plot_tree
import matplotlib.pyplot as plt

# Determine how many stumps (weak Learners)
n_stumps = len(ada.estimators_)

# Create subplots dynamically
fig, axes = plt.subplots(1, n_stumps, figsize=(5 * n_stumps, 4))

# Handle case when only one stump exists
```

```

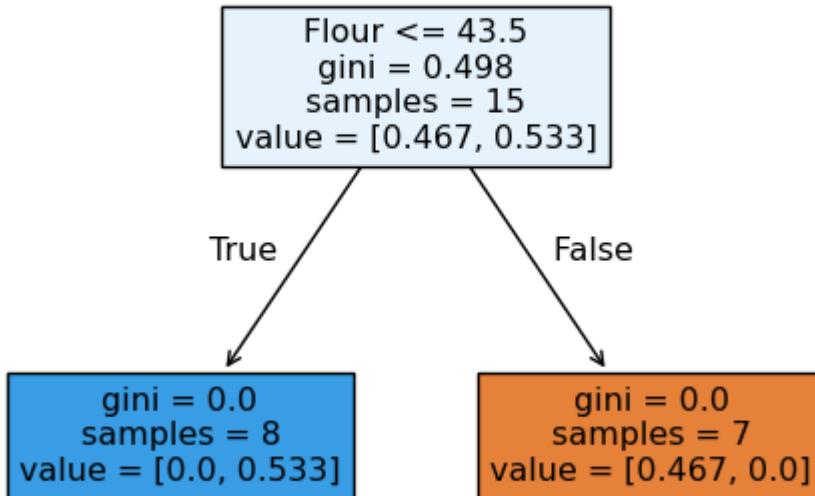
if n_stumps == 1:
    axes = [axes]

# Plot each decision stump on its corresponding axis
for i, tree in enumerate(ada.estimators_):
    plot_tree(tree, feature_names=['Sugar', 'Flour'], filled=True, ax=axes[i])
    axes[i].set_title(f"Decision Stump {i + 1}")

plt.tight_layout()
plt.show()

```

## Decision Stump 1



In [36]:

```
# 6 Plot final decision boundary
x_min, x_max = X['Sugar'].min() - 1, X['Sugar'].max() + 1
y_min, y_max = X['Flour'].min() - 1, X['Flour'].max() + 1
xx, yy = np.meshgrid(np.linspace(x_min, x_max, 200),
                      np.linspace(y_min, y_max, 200))
```

In [38]:

```
Z = ada.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)
```

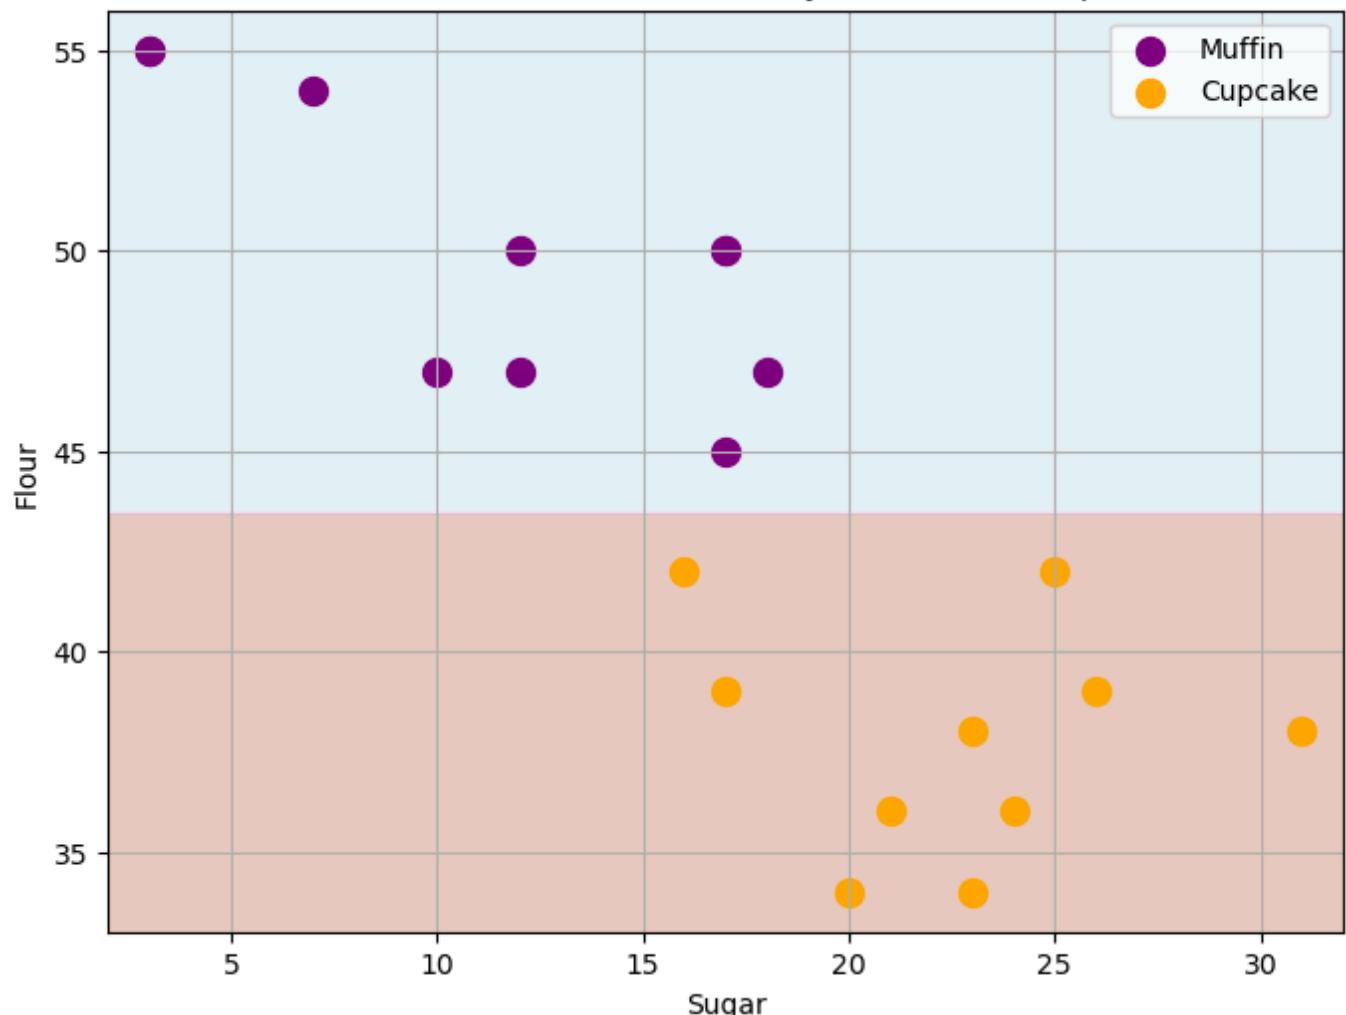
C:\Users\Admin\.conda\envs\myenv\Lib\site-packages\sklearn\utils\validation.py:2749: UserWarning: X does not have valid feature names, but AdaBoostClassifier was fitted with feature names  
warnings.warn(

In [39]:

```
plt.figure(figsize=(8,6))
plt.contourf(xx, yy, Z, alpha=0.3, cmap=plt.cm.Paired)
plt.scatter(df[df['Type']=='Muffin']['Sugar'],
            df[df['Type']=='Muffin']['Flour'],
            color='purple', label='Muffin', s=100)
plt.scatter(df[df['Type']=='Cupcake']['Sugar'],
            df[df['Type']=='Cupcake']['Flour'],
            color='orange', label='Cupcake', s=100)

plt.xlabel('Sugar')
plt.ylabel('Flour')
plt.title('Final AdaBoost Decision Boundary - Muffins vs Cupcakes')
plt.legend()
plt.grid(True)
plt.show()
```

### Final AdaBoost Decision Boundary - Muffins vs Cupcakes



In [ ]: