

```
In [81]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.tree import DecisionTreeClassifier, plot_tree
from mlxtend.plotting import plot_decision_regions
from sklearn.preprocessing import LabelEncoder
import warnings
warnings.filterwarnings("ignore", category=pd.errors.SettingWithCopyWarning)
```

```
In [82]: df = pd.read_csv(r"C:\Users\Admin\Downloads\cupcakes - cupcakes.csv") # change path as needed
df.head()
```

```
Out[82]:
```

	Type	Flour	Milk	Sugar	Butter	Egg	Baking Powder	Vanilla	Salt	
0	Muffin	55	28	3	7	5		2	0	0
1	Muffin	47	24	12	6	9		1	0	0
2	Muffin	47	23	18	6	4		1	0	0
3	Muffin	45	11	17	17	8		1	0	0
4	Muffin	50	25	12	6	5		2	1	0

```
In [83]: df.shape
```

```
Out[83]: (20, 9)
```

```
In [84]: le = LabelEncoder()
df['label'] = le.fit_transform(df['Type'])
df['label']
```

```
Out[84]:
```

0	1
1	1
2	1
3	1
4	1
5	1
6	1
7	1
8	1
9	1
10	0
11	0
12	0
13	0
14	0
15	0
16	0
17	0
18	0
19	0


Name: label, dtype: int64

```
In [85]: df = df[["Flour", "Sugar", "label"]]
df.columns = ["X1", "X2", "label"]
```

```
In [86]: df["weights"] = 1 / df.shape[0]
display(df.head())
```

	X1	X2	label	weights
0	55	3	1	0.05
1	47	12	1	0.05
2	47	18	1	0.05
3	45	17	1	0.05
4	50	12	1	0.05

```
In [87]: sns.scatterplot(x="X1", y="X2", hue="label", data=df, s=60)
plt.title("Initial Cupcakes Data (2-class)")
plt.xlabel("Flour")
plt.ylabel("Sugar")
plt.show()
```



```
In [88]: x = df[["X1", "X2"]].values
y = df["label"].values
```

```
In [89]: def calculate_model_weight(error):
    eps = 1e-10
    error = np.clip(error, eps, 1 - eps)
    return 0.5 * np.log((1 - error) / error)
```

```
In [90]: def update_row_weights(row, alpha):
    if row["label"] == row["y_pred"]:
        return row["weights"] * np.exp(-alpha)
    else:
        return row["weights"] * np.exp(alpha)

def create_new_dataset(df):
    indices = []
    n = df.shape[0]
```

```

for _ in range(n):
    a = np.random.random()
    for idx, row in df.iterrows():
        if row["cumsum_lower"] < a <= row["cumsum_upper"]:
            indices.append(idx)
            break
    if len(indices) == 0:
        indices = np.random.choice(df.index, size=n, replace=True).tolist()
return indices

```

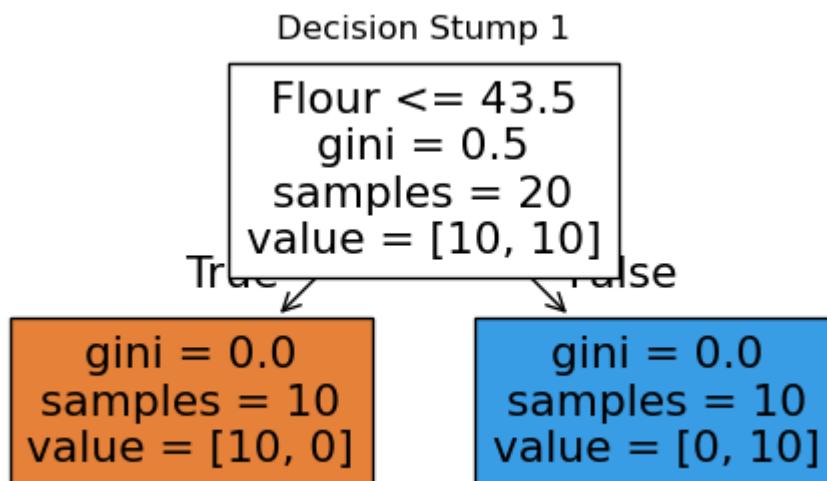
In [91]: #Step 2: Train First Decision Stump
`dt1 = DecisionTreeClassifier(max_depth=1)
dt1.fit(x, y)`

```

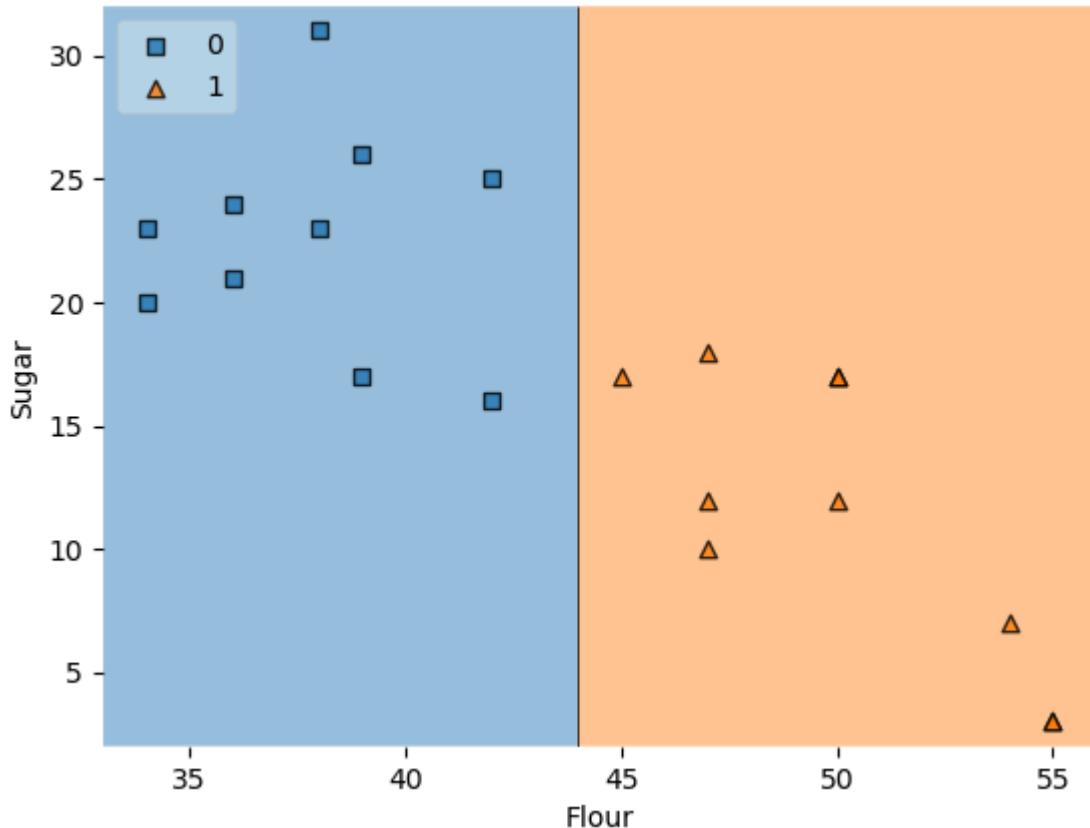
plt.figure(figsize=(6,3))
plot_tree(dt1, filled=True, feature_names=["Flour", "Sugar"])
plt.title("Decision Stump 1")
plt.show()

plot_decision_regions(x, y, clf=dt1, legend=2)
plt.title("Decision Region - Stump 1")
plt.xlabel("Flour")
plt.ylabel("Sugar")
plt.show()

```



Decision Region - Stump 1



```
In [92]: # Predictions & weighted error
df["y_pred"] = dt1.predict(x)
error = np.sum(df["weights"] * (df["label"] != df["y_pred"]))
alpha1 = calculate_model_weight(error)
print(f"Model 1 alpha: {alpha1:.3f}")

# Update and normalize weights
df["updated_weights"] = df.apply(lambda row: update_row_weights(row, alpha1), axis=1)
df["normalized_weights"] = df["updated_weights"] / df["updated_weights"].sum()

# Compute cumulative bounds
df["cumsum_upper"] = np.cumsum(df["normalized_weights"])
df["cumsum_lower"] = df["cumsum_upper"] - df["normalized_weights"]
display(df[["X1", "X2", "label", "weights", "y_pred", "normalized_weights"]])
```

Model 1 alpha: 11.513

	X1	X2	label	weights	y_pred	normalized_weights
0	55	3	1	0.05	1	0.05
1	47	12	1	0.05	1	0.05
2	47	18	1	0.05	1	0.05
3	45	17	1	0.05	1	0.05
4	50	12	1	0.05	1	0.05
5	55	3	1	0.05	1	0.05
6	54	7	1	0.05	1	0.05
7	47	10	1	0.05	1	0.05
8	50	17	1	0.05	1	0.05
9	50	17	1	0.05	1	0.05
10	39	26	0	0.05	0	0.05
11	42	16	0	0.05	0	0.05
12	34	20	0	0.05	0	0.05
13	39	17	0	0.05	0	0.05
14	38	23	0	0.05	0	0.05
15	42	25	0	0.05	0	0.05
16	36	21	0	0.05	0	0.05
17	38	31	0	0.05	0	0.05
18	36	24	0	0.05	0	0.05
19	34	23	0	0.05	0	0.05

```
In [93]: # Resample for next stump
index_values = create_new_dataset(df)
print("Resampled indices (1st):", index_values)
second_df = df.loc[index_values, ["X1", "X2", "label", "normalized_weights"]].reset_index(drop=True)
```



Resampled indices (1st): [13, 17, 11, 10, 17, 17, 18, 9, 3, 4, 14, 6, 7, 17, 2, 17, 16, 5, 7, 13]

```
In [94]: # =====
# Step 3: Train Second Decision Stump
# =====
x2 = second_df[["X1", "X2"]].values
y2 = second_df["label"].values
dt2 = DecisionTreeClassifier(max_depth=1)
dt2.fit(x2, y2)

plt.figure(figsize=(6,3))
plot_tree(dt2, filled=True, feature_names=["Flour", "Sugar"])
plt.title("Decision Stump 2")
plt.show()

plot_decision_regions(x2, y2, clf=dt2, legend=2)
plt.title("Decision Region - Stump 2")
plt.xlabel("Flour")
```

```
plt.ylabel("Sugar")
plt.show()
```



```
In [95]: second_df["y_pred"] = dt2.predict(x2)
error2 = np.sum(second_df["normalized_weights"] * (second_df["label"] != second_df["y_pred"]))
alpha2 = calculate_model_weight(error2)
print(f"Model 2 alpha: {alpha2:.3f}")
```

Model 2 alpha: 11.513

```
In [96]: # Update weights again
def update_row_weights_2(row, alpha=alpha2):
    if row["label"] == row["y_pred"]:
        return row["normalized_weights"] * np.exp(-alpha)
    else:
        return row["normalized_weights"] * np.exp(alpha)

second_df["updated_weights"] = second_df.apply(update_row_weights_2, axis=1)
second_df["normalized_weights"] = second_df["updated_weights"] / second_df["updated_weights"]
second_df["cumsum_upper"] = np.cumsum(second_df["normalized_weights"])
```

```

second_df["cumsum_lower"] = second_df["cumsum_upper"] - second_df["normalized_weights"]

index_values2 = create_new_dataset(second_df)
print("Resampled indices (2nd):", index_values2)
third_df = second_df.loc[index_values2, ["X1", "X2", "label", "normalized_weights"]].reset_index()

```

Resampled indices (2nd): [16, 6, 7, 3, 12, 10, 19, 5, 3, 8, 4, 5, 11, 17, 13, 3, 5, 16, 19, 13]


```

In [97]: # =====
# Step 4: Train Third Decision Stump
# =====
x3 = third_df[["X1", "X2"]].values
y3 = third_df["label"].values
dt3 = DecisionTreeClassifier(max_depth=1)
dt3.fit(x3, y3)


plt.figure(figsize=(6,3))
plot_tree(dt3, filled=True, feature_names=["Flour", "Sugar"])
plt.title("Decision Stump 3")
plt.show()

plot_decision_regions(x3, y3, clf=dt3, legend=2)
plt.title("Decision Region - Stump 3")
plt.xlabel("Flour")
plt.ylabel("Sugar")
plt.show()

```



Decision Region - Stump 3



```
In [98]: third_df["y_pred"] = dt3.predict(x3)
error3 = np.sum(third_df["normalized_weights"] * (third_df["label"] != third_df["y_pred"]))
alpha3 = calculate_model_weight(error3)
print(f"Model 3 alpha: {alpha3:.3f}")

Model 3 alpha: 11.513
```

```
In [99]: # Step 5: Combined Visualization
# =====
print("\nFinal Model Weights (alphas):")
print(f"Alpha1 = {alpha1:.3f}, Alpha2 = {alpha2:.3f}, Alpha3 = {alpha3:.3f}")


plt.figure(figsize=(12,5))

plt.subplot(1,2,1)
plot_decision_regions(x, y, clf=dt1, legend=0)
plot_decision_regions(x, y, clf=dt2, legend=0)
plot_decision_regions(x, y, clf=dt3, legend=0)
plt.title("Decision Regions: All 3 Stumps")
plt.xlabel("Flour")
plt.ylabel("Sugar")
```

Final Model Weights (alphas):
Alpha1 = 11.513, Alpha2 = 11.513, Alpha3 = 11.513

Out[99]: Text(0, 0.5, 'Sugar')

Decision Regions: All 3 Stumps

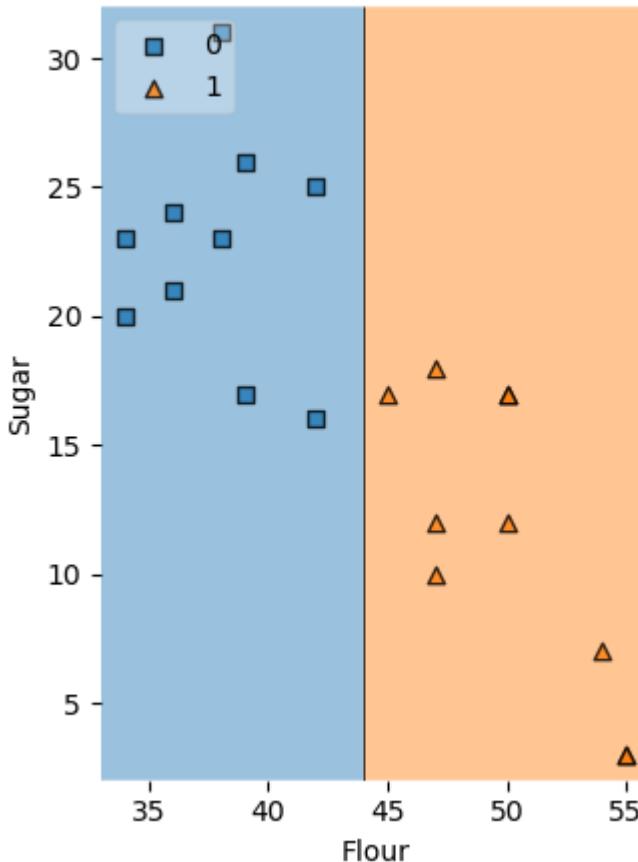


```
In [100]: from sklearn.ensemble import AdaBoostClassifier
ada_final = AdaBoostClassifier(
    estimator=DecisionTreeClassifier(max_depth=1),
    n_estimators=3,
    algorithm="SAMME")
)
ada_final.fit(x, y)

plt.subplot(1,2,2)
plot_decision_regions(x, y, clf=ada_final, legend=2)
plt.title("Final AdaBoost Combined Decision Region")
plt.xlabel("Flour")
plt.ylabel("Sugar")
plt.tight_layout()
plt.show()
```

C:\ProgramData\anaconda3\Lib\site-packages\sklearn\ensemble_weight_boosting.py:519: FutureWarning: The parameter 'algorithm' is deprecated in 1.6 and has no effect. It will be removed in version 1.8.
warnings.warn(

Final AdaBoost Combined Decision Region



In [101]:

```
# =====
# Step 6: Visualize All Stumps Together
# =====
from sklearn.ensemble import AdaBoostClassifier

plt.figure(figsize=(14, 6))

# --- Subplot 1: Decision Stump 1 ---
plt.subplot(1, 4, 1)
plot_decision_regions(x, y, clf=dt1, legend=0)
plt.title("Decision Stump 1")
plt.xlabel("X1")
plt.ylabel("X2")

# --- Subplot 2: Decision Stump 2 ---
plt.subplot(1, 4, 2)
plot_decision_regions(x2, y2, clf=dt2, legend=0)
plt.title("Decision Stump 2")
plt.xlabel("X1")
plt.ylabel("X2")

# --- Subplot 3: Decision Stump 3 ---
plt.subplot(1, 4, 3)
plot_decision_regions(x3, y3, clf=dt3, legend=0)
plt.title("Decision Stump 3")
plt.xlabel("X1")
plt.ylabel("X2")

# --- Subplot 4: Final Combined AdaBoost ---
ada_final = AdaBoostClassifier(
    estimator=DecisionTreeClassifier(max_depth=1),
    n_estimators=3,
    algorithm="SAMME"
)
ada_final.fit(x, y)
```


```

plt.subplot(1, 4, 4)
plot_decision_regions(x, y, clf=ada_final, legend=2)
plt.title("Final AdaBoost Combined")
plt.xlabel("X1")
plt.ylabel("X2")

plt.tight_layout()
plt.show()

```

C:\ProgramData\anaconda3\Lib\site-packages\sklearn\ensemble_weight_boosting.py:519: FutureWarning: The parameter 'algorithm' is deprecated in 1.6 and has no effect. It will be removed in version 1.8.



In [102...]

```

# -----
# 🍪 Muffin vs Cupcake using AdaBoost + ALL Decision Stumps Visualization
# -----
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.ensemble import AdaBoostClassifier

```

In [103...]

```
df = pd.read_csv(r"C:\Users\Admin\Downloads\cupcakes - cupcakes.csv") # change path as needed
df.head()
```

Out[103...]

	Type	Flour	Milk	Sugar	Butter	Egg	Baking Powder	Vanilla	Salt
0	Muffin	55	28	3	7	5	2	0	0
1	Muffin	47	24	12	6	9	1	0	0
2	Muffin	47	23	18	6	4	1	0	0
3	Muffin	45	11	17	17	8	1	0	0
4	Muffin	50	25	12	6	5	2	1	0

In [104...]

```

le = LabelEncoder()
df['label'] = le.fit_transform(df['Type'])
df['label']

```

```
Out[104...]: 0      1  
1      1  
2      1  
3      1  
4      1  
5      1  
6      1  
7      1  
8      1  
9      1  
10     0  
11     0  
12     0  
13     0  
14     0  
15     0  
16     0  
17     0  
18     0  
19     0  
Name: label, dtype: int64
```

```
In [105...]: df['Label'] = df['Type'].map({'Muffin': 0, 'Cupcake': 1})
```

```
In [106...]: X = df[['Sugar', 'Flour']]  
y = df['Label']
```

```
In [107...]: # 2 Split into training and test data  
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42)
```

```
In [108...]: # 3 Create AdaBoost model with Decision Stumps  
from sklearn.ensemble import AdaBoostClassifier  
from sklearn.tree import DecisionTreeClassifier  
base_tree = DecisionTreeClassifier(max_depth=1)  
ada = AdaBoostClassifier(  
    estimator=base_tree,  
    n_estimators=3,          # use 3 stumps for visualization  
    learning_rate=1.0,  
    random_state=42  
)
```

```
In [109...]: # 4 Train model  
ada.fit(X_train, y_train)
```

```
Out[109...]: ► AdaBoostClassifier  
             (i) (?)  
             ► estimator:  
                 DecisionTreeClassifier  
                     ► DecisionTreeClassifier (i)
```

```
In [110...]: from sklearn.tree import plot_tree  
import matplotlib.pyplot as plt  
  
# Determine how many stumps (weak Learners)  
n_stumps = len(ada.estimators_)  
  
# Create subplots dynamically  
fig, axes = plt.subplots(1, n_stumps, figsize=(5 * n_stumps, 4))
```

```


# Handle case when only one stump exists
if n_stumps == 1:
    axes = [axes]

# Plot each decision stump on its corresponding axis
for i, tree in enumerate(ada.estimators_):
    plot_tree(tree, feature_names=['Sugar', 'Flour'], filled=True, ax=axes[i])
    axes[i].set_title(f"Decision Stump {i + 1}")

plt.tight_layout()
plt.show()

```

Decision Stump 1



```
In [111]: # 6 Plot final decision boundary
x_min, x_max = X['Sugar'].min() - 1, X['Sugar'].max() + 1
y_min, y_max = X['Flour'].min() - 1, X['Flour'].max() + 1
xx, yy = np.meshgrid(np.linspace(x_min, x_max, 200),
                      np.linspace(y_min, y_max, 200))
```

```
In [112]: Z = ada.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)
```


C:\ProgramData\anaconda3\Lib\site-packages\sklearn\utils\validation.py:2739: UserWarning: X does not have valid feature names, but AdaBoostClassifier was fitted with feature names
 warnings.warn(

```
In [113]: plt.figure(figsize=(8,6))
plt.contourf(xx, yy, Z, alpha=0.3, cmap=plt.cm.Paired)
plt.scatter(df[df['Type']=='Muffin']['Sugar'],
            df[df['Type']=='Muffin']['Flour'],
            color='purple', label='Muffin', s=100)
plt.scatter(df[df['Type']=='Cupcake']['Sugar'],
            df[df['Type']=='Cupcake']['Flour'],
            color='orange', label='Cupcake', s=100)

plt.xlabel('Sugar')
plt.ylabel('Flour')
plt.title('Final AdaBoost Decision Boundary - Muffins vs Cupcakes')
plt.legend()
```

```
plt.grid(True)  
plt.show()
```

Final AdaBoost Decision Boundary - Muffins vs Cupcakes



In [114]:

```
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
from sklearn.tree import DecisionTreeClassifier  
from sklearn.ensemble import AdaBoostClassifier  
from mlxtend.plotting import plot_decision_regions
```

In [115]:

```
data = pd.read_csv(r"C:\Users\Admin\Downloads\cupcakes - cupcakes.csv") # change path as needed  
data.head()
```

Out[115]:

	Type	Flour	Milk	Sugar	Butter	Egg	Baking Powder	Vanilla	Salt	
0	Muffin	55	28	3	7	5		2	0	0
1	Muffin	47	24	12	6	9		1	0	0
2	Muffin	47	23	18	6	4		1	0	0
3	Muffin	45	11	17	17	8		1	0	0
4	Muffin	50	25	12	6	5		2	1	0

In [116]:

```
# Convert to numeric labels  
data['Label'] = data['Type'].map({'Muffin': 0, 'Cupcake': 1})  
X = data[['Sugar', 'Flour']].values  
y = data['Label'].values
```

In [117...]

```
base_tree = DecisionTreeClassifier(max_depth=1, random_state=42)

ada = AdaBoostClassifier(
    estimator=base_tree,
    n_estimators=3,
    learning_rate=1.0,
    random_state=42
)
```

In [118...]

```
ada.fit(X, y)
print("Number of trained stumps:", len(ada.estimators_))

plt.figure(figsize=(16, 5))
estimators = ada.estimators_

# Decision Stump 1
plt.subplot(1, 4, 1)
plot_decision_regions(X, y, clf=estimators[0], legend=0)
plt.title("Decision Stump 1")
plt.xlabel("Sugar")
plt.ylabel("Flour")



# Decision Stump 2
if len(estimators) > 1:
    plt.subplot(1, 4, 2)
    plot_decision_regions(X, y, clf=estimators[1], legend=0)
    plt.title("Decision Stump 2")
    plt.xlabel("Sugar")
    plt.ylabel("Flour")

# Decision Stump 3
if len(estimators) > 2:
    plt.subplot(1, 4, 3)
    plot_decision_regions(X, y, clf=estimators[2], legend=0)
    plt.title("Decision Stump 3")
    plt.xlabel("Sugar")
    plt.ylabel("Flour")

# Final Combined AdaBoost
plt.subplot(1, 4, 4)
plot_decision_regions(X, y, clf=ada, legend=2)
plt.title("Final AdaBoost Combined")
plt.xlabel("Sugar")
plt.ylabel("Flour")

plt.tight_layout()
plt.show()
```

Number of trained stumps: 1



```
In [119]: print("\nAlpha (Model Weights):", ada.estimator_weights_)
print("Errors of stumps:", ada.estimator_errors_)
```

Alpha (Model Weights): [1. 0. 0.]

Errors of stumps: [0. 1. 1.]

```
In [121]: sample = np.array([[22, 46]]) # (Sugar, Flour)
pred = ada.predict(sample)[0]
label = "Cupcake" if pred == 1 else "Muffin"
print(f"\nNew recipe with Sugar={sample[0][0]}, Flour={sample[0][1]} → Predicted: {label}")
```

New recipe with Sugar=22, Flour=46 → Predicted: Muffin